

Deploying ML Models on AWS platform

Agenda

- ❖ Train a model locally. (Local environment should be ready).
- ❖ Upload the model to AWS S3 bucket.
- ❖ Create a lambda function.
- ❖ Integrate with API gateway.
- ❖ Optional (discussion on layers)
- ❖ Optional (discussion on how to productionize DNN models on Lambda)

INTRODUCTIONS

A Request

- ❖ This will be a hands-on session
- ❖ Please do not ask code for copy paste. This exercise to make you implement a ml model, end to end, copy paste will not serve the purpose.
- ❖ Please remain actively engaged during workshop and follow facilitator's guidance.
- ❖ Please keep your video on for the entire session duration
- ❖ Feedback link available below must be filled. That will be also the record of your participation in this session.
- ❖ Do ask for breaks.

Features of AWS Lambda

- ❖ For hosting small models
- ❖ Serverless
- ❖ Suitable for Low-cost scenario, experimentation. When you need to have a end to end demo.
- ❖ No GPU
- ❖ Deep Neural Networks of reasonable size can be productionized.
- ❖ Suitable for domains like Manufacturing.
- ❖ No heavy infrastructure, less maintenance.
- ❖ ML Pipeline can be established (retraining) – Can address some of the ML lifecycle challenges.

Lambda Limits

Resource	Quota
Function memory allocation	128 MB to 10,240 MB, in 1-MB increments.
Function timeout	900 seconds (15 minutes)
Function environment variables	4 KB
Function resource-based policy	20 KB
Function layers	five layers
Function burst concurrency	500 - 3000 (varies per Region)
Invocation payload (request and response)	6 MB (synchronous)
	256 KB (asynchronous)
Deployment package (.zip file archive) size	50 MB (zipped, for direct upload)

Resource	Quota
	256 KB (asynchronous)
Deployment package (.zip file archive) size	50 MB (zipped, for direct upload)
	250 MB (unzipped, including layers)
	3 MB (console editor), 512 KB maximum for an individual file
Container image code package size	10 GB
Test events (console editor)	10
/tmp directory storage	512 MB
File descriptors	1,024
Execution processes/threads	1,024

Part – A - Train the model (if we do not have)

- ❖ Use Wine dataset for training a model.
- ❖ Save the model to disk.
- ❖ Load the model from the disk.
- ❖ Do prediction on sample data.

Train the model

```
from sklearn.datasets import load_wine
import pandas as pd
import numpy as np
import pickle

data = load_wine() # import dataset
df = pd.DataFrame(data['data'], columns=data['feature_names']) # build dataframe
df['target'] = data['target'] # add dependent variable
df = df.sample(frac=1) # randomize the data
df.head(3)

train_df = df[:150]
test_df = df[150:]

feature_cols = ['alcohol', 'malic_acid']
X = train_df[feature_cols] # Features
y = train_df.target # Target variable

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=0)
```


Train the model

```
from sklearn.linear_model import LogisticRegression

# instantiate the model (using the default parameters)
logreg = LogisticRegression()

# fit the model with data
logreg.fit(X_train,y_train)
y_pred=logreg.predict(X_test)
```

Testing (Predicting) for one observation:

```
data = {
    "alcohol": 11.1, "malic_acid": 5
}
df = pd.DataFrame([data])
result = logreg.predict(df)
result[0]
```

Creating Model file

```
import pickle  
pickle.dump( logreg, open( "wine-model-masterclass-2-input-params.p", "wb" ))
```

change the name so that it is unique to you <model-wine-ritesh.p>

Testing the model from disk at Training time:

- This is an important step.
- This is setting up the code which can be used at the time of setting this up.
- if you have done this step well, rest of things will be easy.
- Load_fn
- Input_fn
- Preprocessing_fn
- Predict_fn
- Output_fn




```
temp_file_path = "wine-model-masterclass-2-input-params.p"
with open(temp_file_path, 'rb') as f:
    model = pickle.load(f)
```

```
input_ = {"alcohol": 105, "malic_acid": 13}
df_input = pd.DataFrame([input_])
df_input
model.predict(df_input)
```

PART B - Build the Lambda function

- ❖ Putting the model on S3 bucket
- ❖ Add Layers (Prebuilt)
- ❖ Add Environment Variables
- ❖ Building Lambda function (Step by Step)

Upload model to AWS

Buckets (24)   Copy ARN  Empty

Buckets are containers for data stored in S3. [Learn more](#)

	Name ▲	AWS Region ▼	Access ▼
			public
<input type="radio"/>	sagemaker-studio-096374906812-rtryv7fu8l	US East (Ohio) us-east-2	Objects can be public
<input type="radio"/>	sagemaker-us-east-2-096374906812	US East (Ohio) us-east-2	Bucket and objects not public
<input type="radio"/>	test-model-bucket	US East (Ohio) us-east-2	Bucket and objects not public
<input type="radio"/>	tfg-models	US East (Ohio) us-east-2	

- Login to AWS account with your credentials.
- Goto S3 Services
- Locate **tfg-models**
- Upload the model file created above.**

Objects (5)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects you'll need to explicitly grant them permissions. [Learn more](#)



List versions



Delete

Actions ▼

Create folder

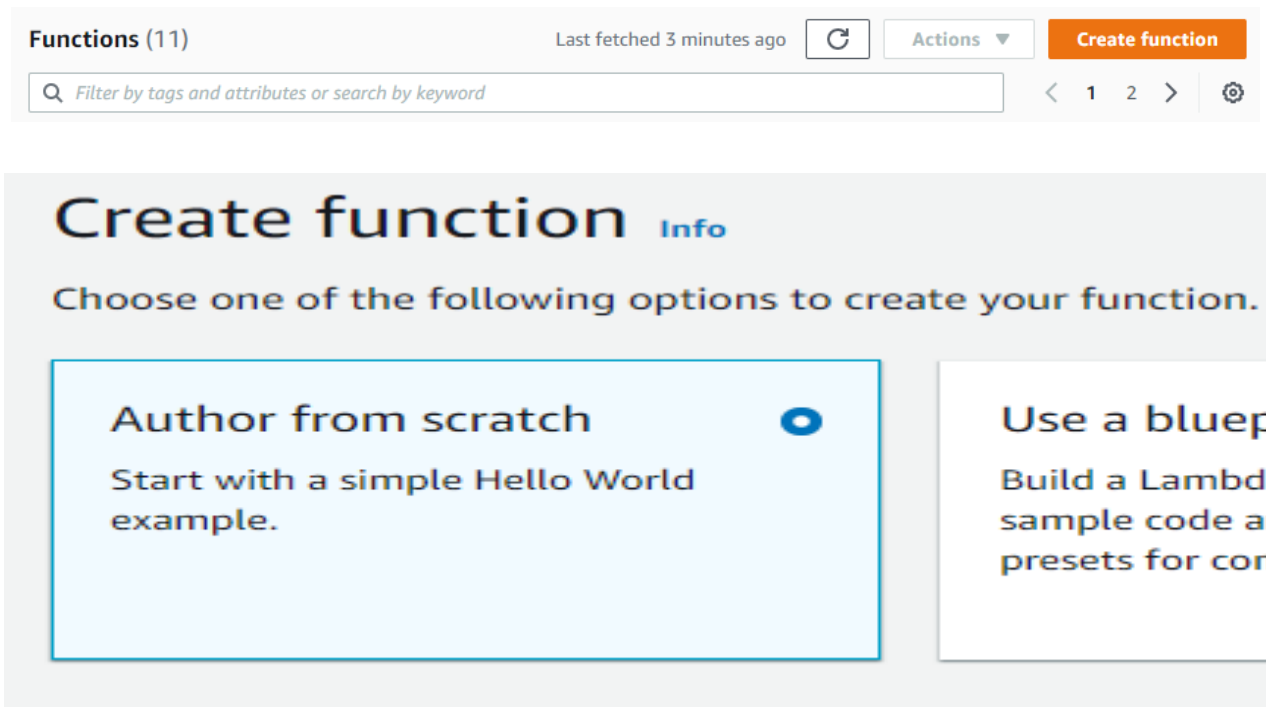
Upload

Building lambda function



Go to Lambda services



<https://us-east-2.console.aws.amazon.com/lambda/home?region=us-east-2#/functions>

Create Function




The screenshot shows the AWS Lambda console interface. At the top, there's a header for 'Functions (11)' with a refresh button, an 'Actions' dropdown, and a 'Create function' button. Below this is a search bar with the placeholder text 'Filter by tags and attributes or search by keyword'. The main content area is titled 'Create function' with an 'Info' link. It instructs the user to 'Choose one of the following options to create your function.' There are two visible options: 'Author from scratch' (selected with a blue radio button) and 'Use a blueprint'. The 'Author from scratch' option includes the text 'Start with a simple Hello World example.' The 'Use a blueprint' option includes the text 'Build a Lambda sample code and presets for com'.

Functions (11) Last fetched 3 minutes ago  Actions  Create function

 Filter by tags and attributes or search by keyword < 1 2 > 

Create function [Info](#)

Choose one of the following options to create your function.

Author from scratch 

Start with a simple Hello World example.

Use a blueprint

Build a Lambda sample code and presets for com

Building lambda function

Basic information

Function name

Enter a name that describes the purpose of your function.

wine-model-ritesh

Use only letters, numbers, hyphens, or underscores with no spaces.

Runtime [Info](#)

Choose the language to use to write your function. Note that the con

Python 3.6

Advanced Settings - Leave as is

▼ Change default execution role

Execution role

Choose a role that defines the permissions of your function. To create a custom role, go to the [IAM console](#).

- ☐ Create a new role with basic Lambda permissions
- ☒ Use an existing role
- ☐ Create a new role from AWS policy templates

Existing role

Choose an existing role that you've created to be used with this Lambda function. The role must have permis

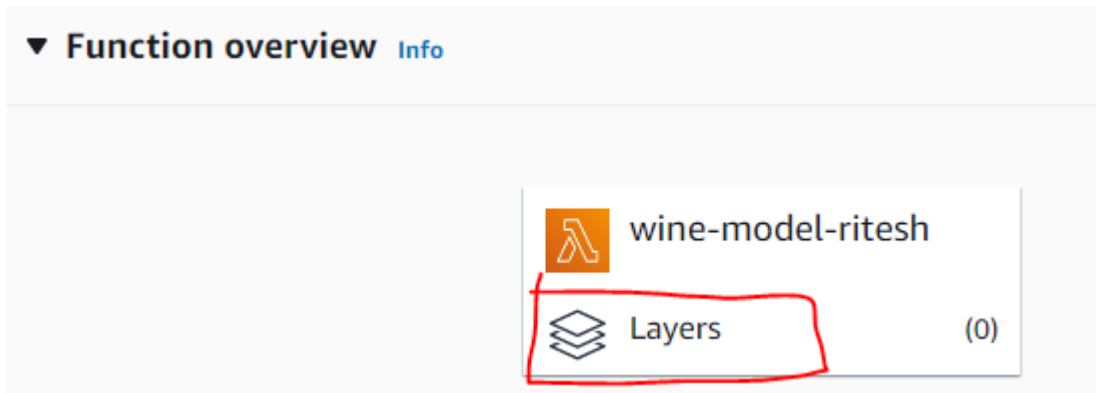
service-role/masterclass-lambda-predict-iris-species-role-pu5jn1j9

[View the masterclass-lambda-predict-iris-species-role-pu5jn1j9 role](#) on the IAM console.

Adding Layers - Building lambda function

Layers provide supporting libraries.

Add following layers:



Layers Info		
Merge order	Name	Layer version
1	AWSLambda-Python36-SciPy1x	35
2	sklearn	8
3	pandas	4

Building lambda function

Version ARN

arn:aws:lambda:us-east-2:259788987135:layer:AWSLambda-Python36-SciPy1x:35

arn:aws:lambda:us-east-2:096374906812:layer:sklearn:8

arn:aws:lambda:us-east-2:096374906812:layer:pandas:4

☒ AWS layers

Choose a layer from a list of layers provided by AWS.

☐ Custom layers

Choose a layer from a list of layers created by your AWS account or organization.

AWS layers

Layers provided by AWS that are compatible with your function's runtime.

AWSLambda-Python36-SciPy1x

Version

35

Choose a layer [Info](#)

Choose from layers with a compatible runtime or specify the Amazon Resource Name (ARN) of

☐ AWS layers

Choose a layer from a list of layers provided by AWS.

☒ Custom layers

Choose a layer from a list of layers created by your AWS account or organization.

Custom layers

Layers created by your AWS account or organization that are compatible with your function's runtime.

sklearn

Version

8

Choose a layer [Info](#)

Choose from layers with a compatible runtime or specify the Amazon Resource Name (ARN) of

☐ AWS layers

Choose a layer from a list of layers provided by AWS.

☒ Custom layers

Choose a layer from a list of layers created by your AWS account or organization.

Custom layers

Layers created by your AWS account or organization that are compatible with your function's runtime.

pandas

Version

4

Define environment variables

Code

Test

Monitor

Configuration

Aliases

Versions

General configuration

Triggers

Permissions

Destinations

Environment variables

Tags

Environment variables (0)

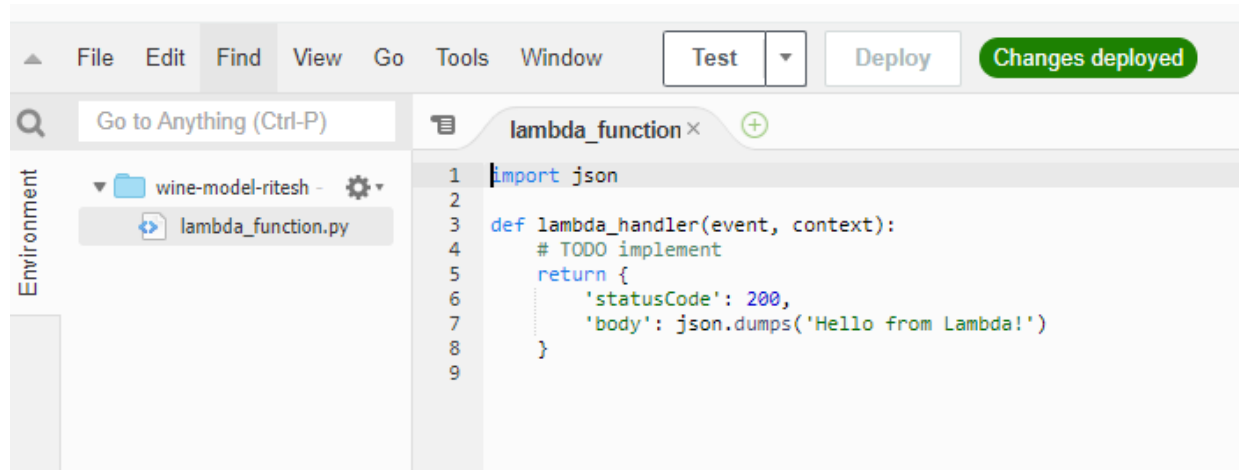
Key	Value
No environment variables	
No environment variables associated with th	
<div>Edit</div>	

MODEL_BUCKET	tfg-models
MODEL_KEY	wine-model-masterclass.p

change

Building lambda function

IMPLEMENT CODE: click on `lambda_function.py`



Building lambda function

Step 1: Test that various parameters and libraries are being imported properly:

```
import json
import sklearn
import boto3
import os
import json
import pickle
import pandas as pd
```

```
s3 = boto3.client('s3')
MODEL_BUCKET=os.environ.get('MODEL_BUCKET')
# get bucket prefix from ENV variable
MODEL_KEY=os.environ.get('MODEL_KEY')
```

```
def lambda_handler(event, context):
    message = f'{MODEL_BUCKET}-{MODEL_KEY}-{sklearn.__version__}'
    # TODO implement
    return {
        'statusCode': 200,
        'body': json.dumps(message)
    }
```

OUTPUT:

Should be something like following:

```
{
  "statusCode": 200,
  "body": "\"tfg-models-wine-model-ritesh.p-0.24.1\""
}
```

Configure Test Event

Configure test event



A function can have up to 10 test events. The events are persisted so you can switch to another computer or web browser and test your function with the same events.

- ☒ Create new test event
- ☐ Edit saved test events

Event template

hello-world



Event name

test

```
1 {  
2   "alcohol": "15",  
3   "malic_acid": "3"  
4 }
```

Loading the model - Building lambda function

MAKE THE FOLLOWING CHANGES TO LOAD THE MODEL

```
s3 = boto3.client('s3')

MODEL_BUCKET=os.environ.get('MODEL_BUCKET')

# get bucket prefix from ENV variable
MODEL_KEY=os.environ.get('MODEL_KEY')

temp_file_path = '/tmp/' + MODEL_KEY

# ** Model Init **
s3.download_file(MODEL_BUCKET, MODEL_KEY, temp_file_path)

print(temp_file_path)

with open(temp_file_path, 'rb') as f:
    model = pickle.load(f)    strclasses = str(model.classes_)
```

```
def lambda_handler(event, context):
    message = f'{MODEL_BUCKET}-{MODEL_KEY}-{strclasses}'
    return {
        'statusCode': 200,
        'body': json.dumps(message)
    }
```

OUTPUT: Response

```
{
  "statusCode": 200,
  "body": "\"tfg-models-wine-model-ritesh.p-[0 1 2]\""
}
```

Prediction - Building lambda function

PREDICTION - Capture Input and Call Predict Function

Inside the lambda handler module:

```
input_ = event
y_pred=model.predict(pd.DataFrame([input_]))
message = f'{MODEL_BUCKET}-{MODEL_KEY}-{y_pred}'
```

Output:

Response

```
{
  "statusCode": 200,
  "body": "\"tfg-models-wine-model-ritesh.p-[2]\""
}
```

Configure test event



A function can have up to 10 test events. The events are persisted so you can switch to another computer or web browser and test your function with the same events.

☐ Create new test event

☒ Edit saved test events

Saved test event

testwinemodel



```
1 {
2   "alcohol": 12,
3   "malic_acid": 15
4 }
```

Building lambda function – Exercise to Participants

Change the lambda function to calculate the time taken in prediction

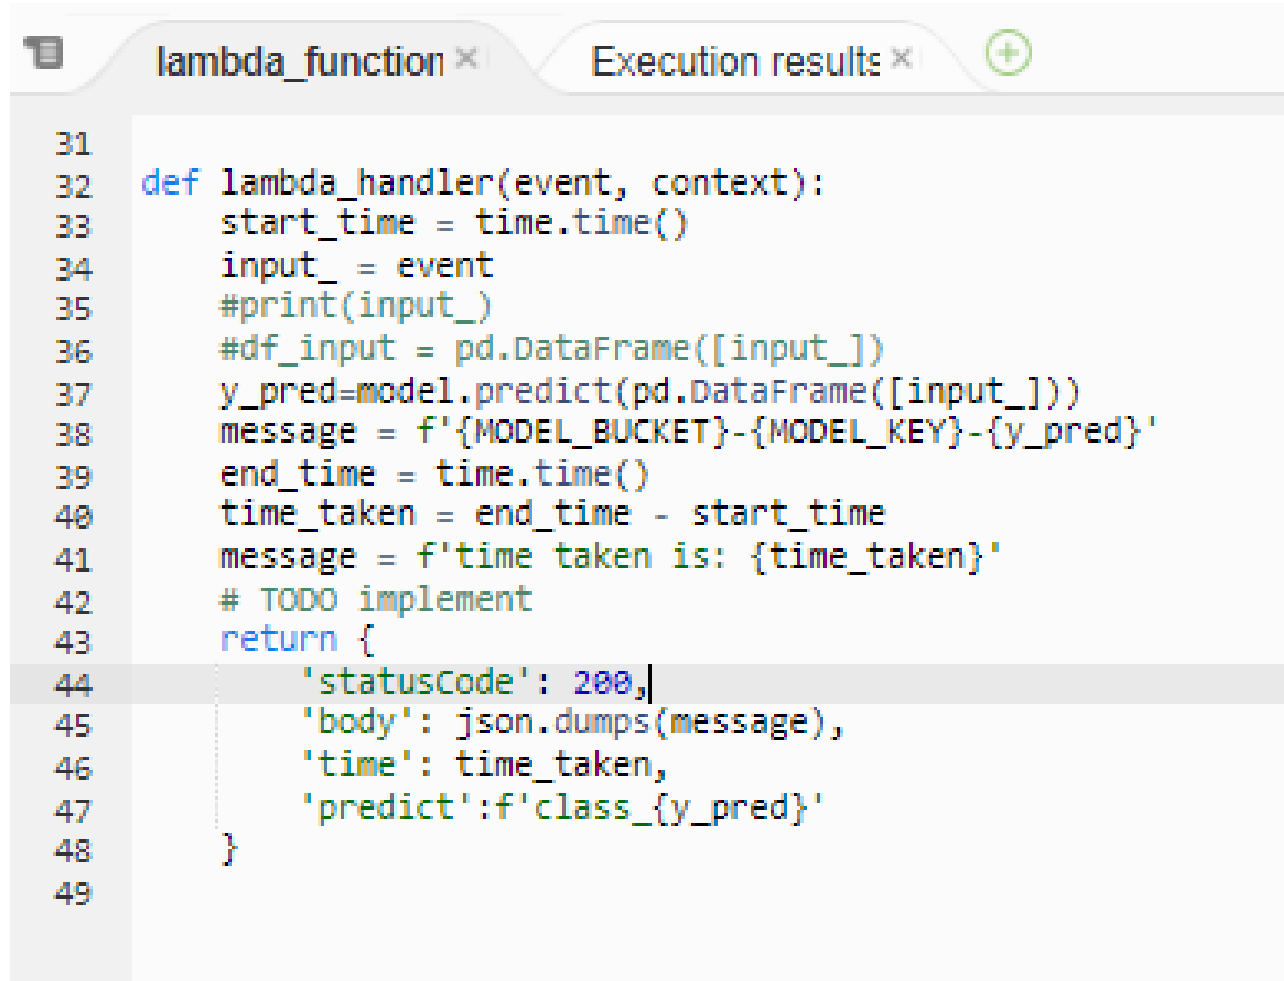
HINT: `start_time = time.time()`

Attach this calculated time taken to the output

Building lambda function – Exercise to Participants

SOLUTION

Building lambda function – Exercise to Participants

A screenshot of a code editor interface. At the top, there are two tabs: 'lambda_function' and 'Execution results'. The 'lambda_function' tab is active, showing a Python code snippet. The code is a lambda function handler that takes an event and context as input. It starts by getting the start time, then gets the input from the event. It prints the input, creates a DataFrame from it, and uses a model to predict a class. It then calculates the end time and the time taken. Finally, it returns a JSON response with the status code 200, the message, the time taken, and the predicted class. The code is as follows:

```
31
32 def lambda_handler(event, context):
33     start_time = time.time()
34     input_ = event
35     #print(input_)
36     #df_input = pd.DataFrame([input_])
37     y_pred=model.predict(pd.DataFrame([input_]))
38     message = f'{MODEL_BUCKET}-{MODEL_KEY}-{y_pred}'
39     end_time = time.time()
40     time_taken = end_time - start_time
41     message = f'time taken is: {time_taken}'
42     # TODO implement
43     return {
44         'statusCode': 200,
45         'body': json.dumps(message),
46         'time': time_taken,
47         'predict':f'class_{y_pred}'
48     }
49
```

Integrate with API gateway

- ❖ Create API
- ❖ Create Resource
- ❖ Create Method
- ❖ Test on Platform
- ❖ Test with Curl

Create API

REST API

Develop a REST API where you gain complete control over the request and response along with API management capabilities.

Works with the following:
Lambda, HTTP, AWS Services

Import

Build

Settings

Choose a friendly name and description for your API.

API name*

test-api

Description

test api

Endpoint Type

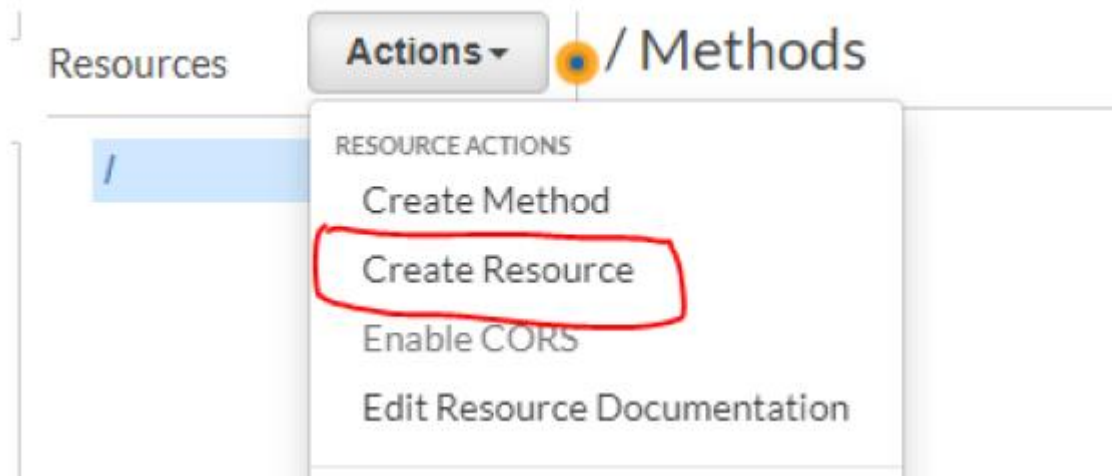
Regional


Create API


In Amazon API Gateway, a REST API refers to a collection of resources and

☒ New API ☐ Clone from existing API ☐ Import

Create Resource




Use this page to create a new child resource for your resource. 

Configure as [proxy resource](#) ☐ 

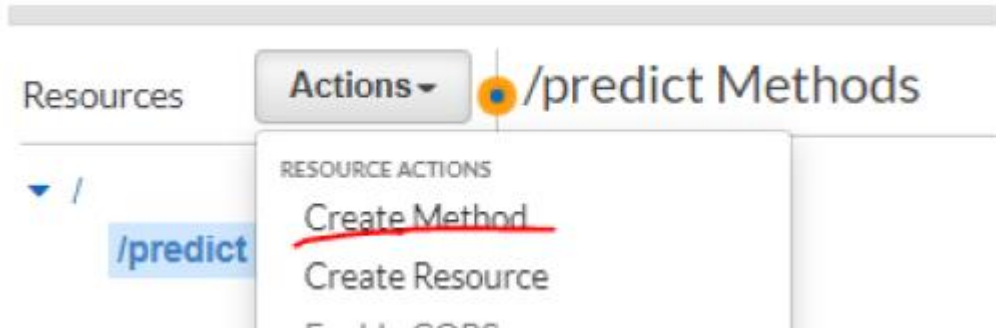
Resource Name*

Resource Path*

You can add path parameters using brackets. For example, you can add a parameter called 'username'. Configuring `/username` as an example, it works for a GET request to `/foo`. To handle

Enable API Gateway CORS ☐ 

Create Method



/predict - POST - Setup

Choose the integration point for your new method.

Integration type ☒ Lambda Function ⓘ
☐ HTTP ⓘ
☐ Mock ⓘ
☐ AWS Service ⓘ
☐ VPC Link ⓘ

Use Lambda Proxy integration ☐ ⓘ

Lambda Region

Lambda Function

Use Default Timeout ☒ ⓘ

Add Permission to Lambda Function

You are about to give API Gateway permission to invoke your Lambda function:
arn:aws:lambda:us-east-2:096374906812:function:wine-model-ritesh

Cancel

OK

Test Method

[← Method Execution](#) /predict - POST - Method Test

Make a test call to your method. When you make a test call, API Gateway skips authorization and directly invokes your method

Path

No path parameters exist for this resource. You can define path parameters by using the syntax {myPathParam} in a resource path.

Query Strings

{predict}

param1=value1¶m2=value2

Headers

{predict}

Use a colon (:) to separate header name and value, and new lines to declare multiple headers. eg.
Accept:application/json.

Stage Variables

No [stage variables](#) exist for this method.

Request Body

```
1 {"alcohol": "12", "malic_acid": "5"}
```

Request: /predict

Status: 200

Latency: 4113 ms

Response Body

```
{
  "statusCode": 200,
  "body": "\"time taken is: 0.4230184555053711\"",
  "time": 0.4230184555053711,
  "predict": "class_[2]"
}
```

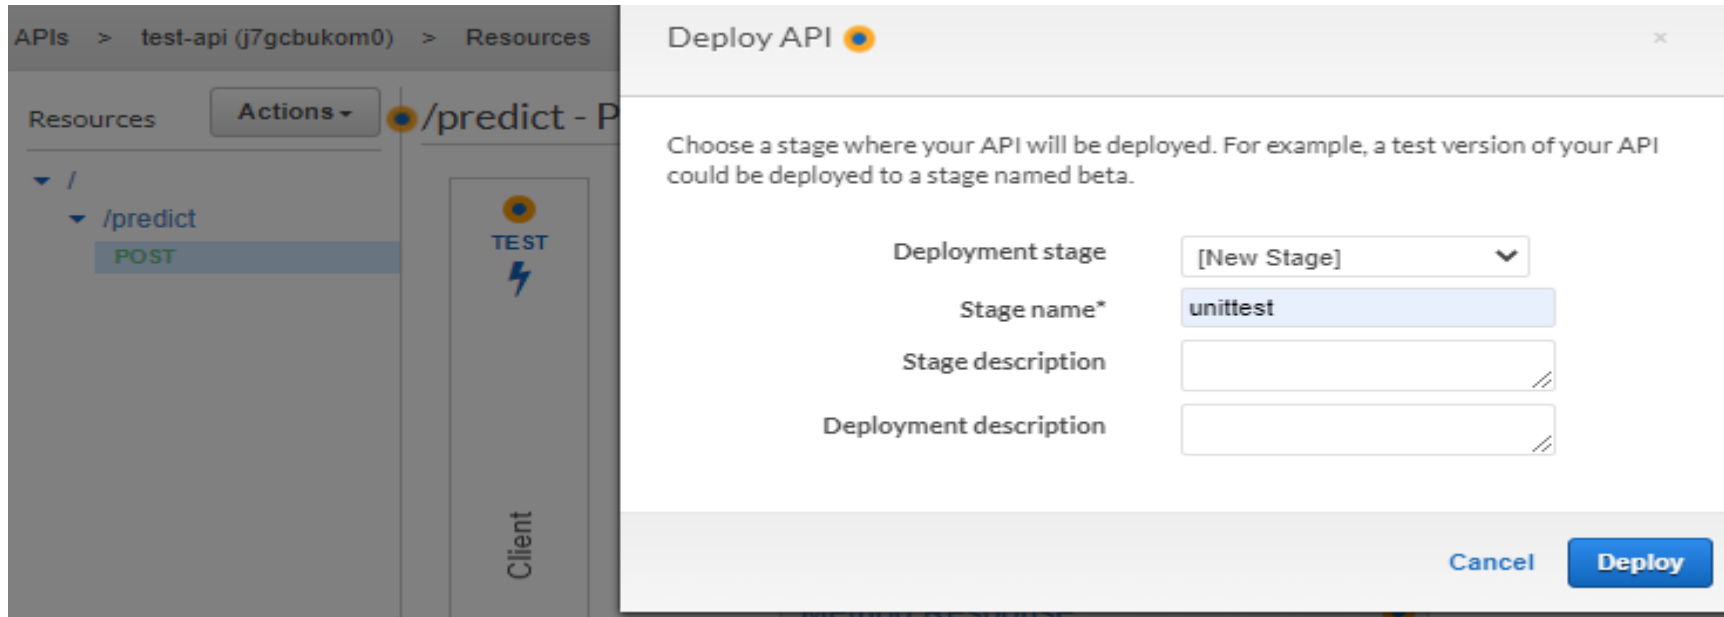
Response Headers

```
{"X-Amzn-Trace-Id":"Root=1-605b2743-e57abf44bbc2983a5l",
"Content-Type":"application/json"}
```

Logs

```
Execution log for request 8613143a-82b4-4bb7-8ed3-66e
Wed Mar 24 11:49:23 UTC 2021 : Starting execution for
8ed3-66e009d65fea
Wed Mar 24 11:49:23 UTC 2021 : HTTP Method: POST, Res
Wed Mar 24 11:49:23 UTC 2021 : Method request path: {
Wed Mar 24 11:49:23 UTC 2021 : Method request query s
Wed Mar 24 11:49:23 UTC 2021 : Method request headers
Wed Mar 24 11:49:23 UTC 2021 : Method request body be
```

Deploy API



Invoke URL: <https://j7gcbukom0.execute-api.us-east-2.amazonaws.com/unittest>

COPY THE URL : <https://j7gcbukom0.execute-api.us-east-2.amazonaws.com/unittest>

Test from command line

```
curl -v -X POST "https://<subdomain>.execute-api.us-east-2.amazonaws.com/unittest/predict/" -H "content-type: application/json" -d "{ \"alcohol\": \"12\", \"malic_acid\": \"15\" }"
```

```
{"statusCode": 200, "body": "\"time taken is: 0.42817234992980957\"", "time": 0.42817234992980957, "predict": "class_[2]"}  
* Connection #0 to host i7chukw0.execute-api.us-east-2.amazonaws.com left intact
```

Discussion

Creating a layer:

```
import os
import shutil
os.listdir('.')
! pip3 install --ignore-installed --target=python pandas
# ! pip3 install --ignore-installed --target=python scikit-learn ' for scikit learn
! rm -rf python/numpy* python/scipy*
! zip -r ./pandas.zip python
```

Deploying a deep NN model

<https://github.com/sinharitesh/lambda-multilabel>

Exercise for participants

Advanced Exercise:

Move the code loading to lambda handler call and note the time taken. is there a drop in subsequent calls. that will tell if there is time difference.

DELETED

BACKUP

Amazon S3 > python-layers-bucket

python-layers-bucket

Objects

Properties

Permissions

Metrics

Management

Access Points

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects you'll need to explicitly grant them permissions. [Learn more](#)

Delete

Actions ▼

Create folder

Upload

< 1 >

<input type="checkbox"/>	Name ▲	Type ▼	Last modified ▼	Size ▼	Storage class
<input type="checkbox"/>	pandas.zip	zip	March 8, 2021, 16:04:13 (UTC+05:30)	14.7 MB	Standard
<input type="checkbox"/>	scikitlearn.zip	zip	March 8, 2021, 12:15:31 (UTC+05:30)	24.4 MB	Standard

Securing the API

APIs > api-wine-test (tfuo2igiwk) > Resources > /predict (uc8g11) > POST

Resources

Actions ▾

← Method Execution

/predict - POST - Method Request

▼ /

▼ /predict


POST


Provide information about this method's authorization settings and the parameters it can receive.

Settings

Authorization


NONE






Request Validator

NONE







API Key Required

true


▼





▶ URL Query String Parameters

▶ HTTP Request Headers

▶ Request Body 

▶ SDK Settings